

# INFO-H503

**Real-time 3D graphics processing and GPGPU imaging  
Projects**

# Overview projects

	Equ/recipes available	C-Code available	CUDA code available
Stereo disparity	Y	Y	N
Light Field depth estimation	Y	N	N
Face detector	Y (in CUDA)	Y	N
Bilateral filtering	Y	Y	Y
Mutual Information Registration	Y	Y	N
2D wavelet transform	Y	Matlab	OpenCL
Cloth simulation	Y	Y/N	N
Sound wave propagation	Y	N	Y/N
3D wavelet transform	Y	N	N
Least square 3D object fitting	Y (in CUDA)	N	N
Surface normal estimation	Y	N	N
Almost rigid deformations	Y	Y	N
Geodesic distances	Y	N	N
Simple ray tracer	Y	Y	Y/N
Volume ray tracer	Y	N	Y

# Attention points

- Understand scientific papers
- Papers are available, sometimes also C/C++ and even CUDA code (in simplified version, as guideline)
- Image processing: in/out = file, real-time rendering is optional
- 3D graphics: minimal OpenGL 3D rendering needed
- Difficult projects: 1-2 students with some modules chosen
- CUDA mandatory, OpenCL optional
- CPU and GPU timing needed
- GPU optimizations needed
- Gradual presentations

# Depth estimation

# Stereo disparity

## Stereo Disparity through Cost Aggregation with Guided Filter

Pauline Tan<sup>1</sup>, Pascal Monasse<sup>2</sup>



(a) Left (reference) image



(b) Right (target) image



```

//*****
//** This method computes the disparity map by means of stereo block matching */
void StereoRig::Compute_Disparity_Map(void)
{
    cvFindStereoCorrespondenceBM(img_left,img_right,disp,BMState);
}


//*****
//** Gets subpixelic disparity for an image position (u,v) of the left image */
double StereoRig::Get_Disparity_32(int u, int v)
{
    double disp_32 = 0.0;

    if( u >= 0 && u < disp->width && v >=0 && v < disp->height )
        disp_32 = cvGetReal2D(disp,v,u)/16.0;



    return disp_32;
}

```

# Light Field Depth Estimation (1-2 students)



Courtesy of UHasselt







# Feature detection and Filtering



# Face detector (1-2 students)

## An Efficient GPGPU Implementation of Viola-Jones Classifier Based Face Detection Algorithm

Vinay Gangadhar      Sharmila Shridhar      Ram Sai Manoj Bamdhamravuri  
 {gangadhar, sshridhar, bamdhamravur}@wisc.edu





# Bilateral filtering



$$GC[I]_{\mathbf{p}} = \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma}(\|\mathbf{p} - \mathbf{q}\|) I_{\mathbf{q}},$$

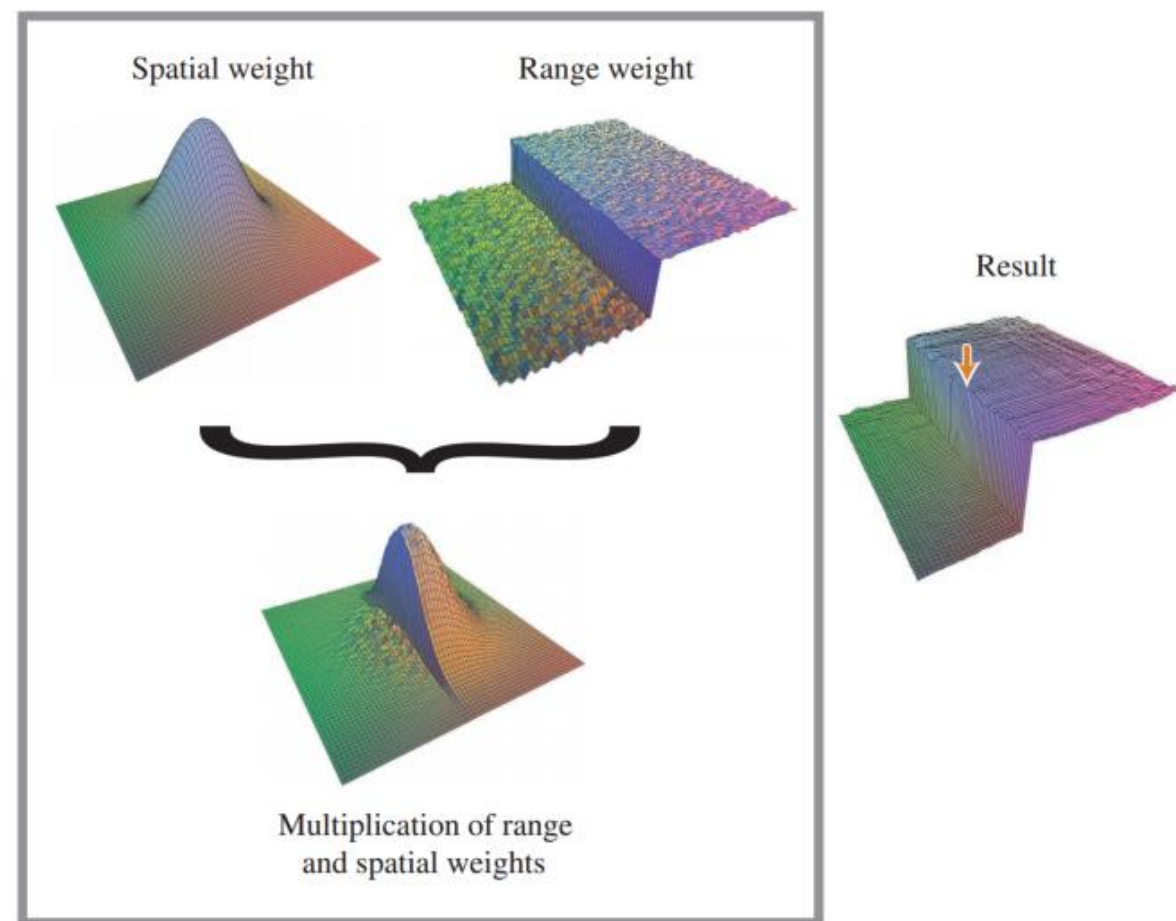

where  $G_{\sigma}(x)$  denotes the 2D Gaussian kernel (see Figure 2.1):

$$G_{\sigma}(x) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right).$$

$$BF[I]_{\mathbf{p}} = \frac{1}{W_{\mathbf{p}}} \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(|I_{\mathbf{p}} - I_{\mathbf{q}}|) I_{\mathbf{q}},$$

where normalization factor  $W_{\mathbf{p}}$  ensures pixel weights sum to 1.0:

$$W_{\mathbf{p}} = \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(|I_{\mathbf{p}} - I_{\mathbf{q}}|).$$





# Bilateral filtering

## Bilateral Filtering with CUDA

Lasse Kløjgaard Staal, 20072300\*


## Efficient and Accurate Gaussian Image Filtering Using Running Sums

Elhanan Elboher and Michael Werman



(a) Gaussian kernel.

(b) SH [24] approximation.



(c) Proposed approximation (4 constants). (d) Proposed approximation (5 constants).

```
__global__
void bilateralFilterGPU_v5(float3* output, uint2 dims, int
radius, float* kernel, float variance, float sqrt_pi_variance)
{
    const unsigned int idx = blockIdx.x*blockDim.x + threadIdx.x;

    uint2 pos = idx_to_co(idx,dims);

    if(pos.x >= dims.x || pos.y >= dims.y) return;

    float3 currentColor = make_float3(tex1Dfetch(tex,
3*idx),tex1Dfetch(tex, 3*idx+1),tex1Dfetch(tex, 3*idx+2));

    float3 res = make_float3(0.0f,0.0f,0.0f);
    float3 normalization = make_float3(0.0f,0.0f,0.0f);
    float3 weight;
    for(int i = -radius; i <= radius; i++) {
        for(int j = -radius; j <= radius; j++) {

            int x_sample = pos.x+i;
            int y_sample = pos.y+j;

            //mirror edges
            if( x_sample < 0) x_sample = -x_sample;
            if( y_sample < 0) y_sample = -y_sample;
            if( x_sample > dims.x - 1) x_sample = dims.x - 1 - i;
            if( y_sample > dims.y - 1) y_sample = dims.y - 1 - j;

            int tempPos =
co_to_idx(make_uint2(x_sample,y_sample),dims);

            float3 tmpColor = make_float3(tex1Dfetch(tex,
3*tempPos),tex1Dfetch(tex, 3*tempPos+1),tex1Dfetch(tex,
3*tempPos+2));

            float gauss_spatial =
kernel[co_to_idx(make_uint2(i+radius,j+radius),make_uint2(radius*
2+1,radius*2+1))];

            weight.x = gauss_spatial *
gaussian1d_gpu_reg((currentColor.x -
tmpColor.x),variance,sqrt_pi_variance);
            weight.y = gauss_spatial *
gaussian1d_gpu_reg((currentColor.y -
tmpColor.y),variance,sqrt_pi_variance);
            weight.z = gauss_spatial *
gaussian1d_gpu_reg((currentColor.z -
tmpColor.z),variance,sqrt_pi_variance);

            normalization = normalization + weight;
            res = res + (tmpColor * weight);

        }
    }

    res.x /= normalization.x;
    res.y /= normalization.y;
    res.z /= normalization.z;
    output[idx] = res;
}
```

```
__global__
void bilateralFilterGPU_v6(float3* output, uint2 dims, int
radius, float* kernel, float variance, float sqrt_pi_variance)
{
    const unsigned int idx = blockIdx.x*blockDim.x + threadIdx.x;

    uint2 pos = idx_to_co(idx,dims);

    if(pos.x >= dims.x || pos.y >= dims.y) return;

    float3 currentColor = make_float3(tex1Dfetch(tex_red,
idx),tex1Dfetch(tex_green, idx),tex1Dfetch(tex_blue, idx));

    float3 res = make_float3(0.0f,0.0f,0.0f);
    float3 normalization = make_float3(0.0f,0.0f,0.0f);

    for(int i = -radius; i <= radius; i++) {
        for(int j = -radius; j <= radius; j++) {

            int x_sample = pos.x+i;
            int y_sample = pos.y+j;

            //mirror edges
            if( x_sample < 0) x_sample = -x_sample;
            if( y_sample < 0) y_sample = -y_sample;
            if( x_sample > dims.x - 1) x_sample = dims.x - 1 - i;
            if( y_sample > dims.y - 1) y_sample = dims.y - 1 - j;

            int tempPos =
co_to_idx(make_uint2(x_sample,y_sample),dims);

            float3 tmpColor = make_float3(tex1Dfetch(tex_red,
tempPos),tex1Dfetch(tex_green, tempPos),tex1Dfetch(tex_blue,
tempPos));

            float gauss_spatial =
kernel[co_to_idx(make_uint2(i+radius,j+radius),make_uint2(radius*
2+1,radius*2+1))];

            float3 gauss_range;
            gauss_range.x = gaussian1d_gpu_reg(currentColor.x -
tmpColor.x, variance, sqrt_pi_variance);
            gauss_range.y = gaussian1d_gpu_reg(currentColor.y -
tmpColor.y, variance, sqrt_pi_variance);
            gauss_range.z = gaussian1d_gpu_reg(currentColor.z -
tmpColor.z, variance, sqrt_pi_variance);

            float3 weight = gauss_spatial * gauss_range;




            normalization = normalization + weight;
            res = res + (tmpColor * weight);

        }
    }

    res.x /= normalization.x;
    res.y /= normalization.y;
    res.z /= normalization.z;
    output[idx] = res;
}
```



# Mutual Information image registration (1-2 students)



Before



After

CUDA code available on <http://users.cecs.anu.edu.au/~ramtin/cuda.htm>

[http://users.cecs.anu.edu.au/~ramtin/papers/2007/DICTA\\_2007a.pdf](http://users.cecs.anu.edu.au/~ramtin/papers/2007/DICTA_2007a.pdf)

<https://lirias.kuleuven.be/bitstream/123456789/28116/1/Maes97TMI.pdf>

# 2D Wavelet Transform

Wavelet transform for video coders  
using GPU

Discrete Wavelet Transform on  
Consumer-Level Graphics Hardware

Tien-Tsin Wong<sup>†</sup> Chi-Sing Leung<sup>‡</sup> Pheng-Ann Heng<sup>†</sup> Jianqing Wang<sup>†</sup>

```
while (yy < 0)
{
    yy = (windowdimY - yy) % (MOD)c windowdimY;
}

// get data from global mem. Two pixels pr. thread
in[ly2][lx2] = INLL(z,yy,xx); /* Scaling s2 */
in[ly2+1][lx2] = INLH(z,yy,xx); /* Scaling s2 */
in[ly2][lx2+1] = INHL(z,yy,xx); /* Scaling d2 */
in[ly2+1][lx2+1] = INHH(z,yy,xx); /* Scaling d2 */


in[ly2][lx2] /= K;
in[ly2+1][lx2] /= K;
in[ly2][lx2+1] *= K;
in[ly2+1][lx2+1] *= K;

/* Compute first pass reconstruction horizontally */
barrier(CLK_LOCAL_MEM_FENCE | CLK_GLOBAL_MEM_FENCE);


in[ly2][lx2+2] -= D * (in[ly2][lx2+1] + in[ly2][lx2+3]); // s1
in[ly2+1][lx2+2] -= D * (in[ly2+1][lx2+1] + in[ly2+1][lx2+3]); // s1

barrier(CLK_LOCAL_MEM_FENCE);


in[ly2][lx2+3] -= C * (in[ly2][lx2+2] + in[ly2][lx2+4]); // d1
in[ly2+1][lx2+3] -= C * (in[ly2+1][lx2+2] + in[ly2+1][lx2+4]); // d1
```



transform rows



transform columns



(a) Mapping to the base positions. (b) Decomposition with boundary extension.

# Volumetric Image & Audio processing



# Sound wave propagation

Jukka Saarelma



Finite-difference time-domain solver for room acoustics using graphics processing units

$$c^2 \nabla^2 p = \frac{\partial^2 p}{\partial t^2},$$

$$c^2 = \kappa \frac{p_0}{\rho_0}.$$

## Sound Wave Propagation Applied in Games

Marcelo Zamith, Erick Passos, Diego Brandão,  
Anselmo Montenegro, Esteban Clua, Mauricio Kischinhevsky, Regina C.P. Leal-Toledo







# Finite Difference Time Domain

## 3D Finite Difference Computation on GPUs using CUDA

Paulius Micikevicius  
NVIDIA  
2701 San Tomas Expressway  
Santa Clara, CA 95050

$$D_{x,y,z}^{t+1} = c_0 D_{x,y,z}^t + \sum_{i=1}^{k/2} c_i \left( D_{x-i,y,z}^t + D_{x+i,y,z}^t + D_{x,y-i,z}^t + D_{x,y+i,z}^t + D_{x,y,z-i}^t + D_{x,y,z+i}^t \right)$$



## Appendix A: CUDA Source Code for a 25-Point Stencil

```
__global__ void fwd_3D_16x16_order8(TYPE *g_output, TYPE *g_input, TYPE *g_vsq, // output initially contains (t-2) step
                                   const int dimx, const int dimy, const int dimz)
{
#define BDIMX    16 // tile (and threadblock) size in x
#define BDIMY    16 // tile (and threadblock) size in y
#define radius    4 // half of the order in space (k/2)

__shared__ float s_data[BDIMY+2*radius][BDIMX+2*radius];

int ix = blockIdx.x*blockDim.x + threadIdx.x;
int iy = blockIdx.y*blockDim.y + threadIdx.y;
int in_idx = iy*dimx + ix; // index for reading input
int out_idx = 0; // index for writing output
int stride = dimx*dimy; // distance between 2D slices (in elements)

float infront1, infront2, infront3, infront4; // variables for input "in front of" the current slice
float behind1, behind2, behind3, behind4; // variables for input "behind" the current slice
float current; // input value in the current slice

int tx = threadIdx.x + radius; // thread's x-index into corresponding shared memory tile (adjusted for halos)
int ty = threadIdx.y + radius; // thread's y-index into corresponding shared memory tile (adjusted for halos)

// fill the "in-front" and "behind" data
behind3 = g_input[in_idx]; in_idx += stride;
behind2 = g_input[in_idx]; in_idx += stride;
behind1 = g_input[in_idx]; in_idx += stride;
current = g_input[in_idx]; out_idx = in_idx; in_idx += stride;
infront1 = g_input[in_idx]; in_idx += stride;
infront2 = g_input[in_idx]; in_idx += stride;
infront3 = g_input[in_idx]; in_idx += stride;
infront4 = g_input[in_idx]; in_idx += stride;




for(int i=radius; i<dimz-radius; i++)
{
    // advance the slice (move the thread-front)
    behind4 = behind3;
    behind3 = behind2;
    behind2 = behind1;
    behind1 = current;
}
```

# 3D graphics processing

# Cloth simulation

## Deformation Constraints in a Mass-Spring Model to Describe Rigid Cloth Behavior



Xavier Provot



# 3D wavelet transform

## Adapted Semi-Regular 3-D Mesh Coding Based on a Wavelet Segmentation

Céline Roudet<sup>a</sup>, Florent Dupont<sup>a</sup> and Atilla Baskurt<sup>b</sup>




```
void ObjFile::computeNewVerts ()
{
    std::vector<int> ks ( newverts_.size (), 0 );
    for ( std::vector<Triangle>::const_iterator t = triangles_.begin(); t != triangles_.end (); ++t )
    {
        for ( int i =0; i < 3; ++i )
        {
            // all vertices are summed double
            newverts_[ t->v[i] ] = newverts_[ t->v[i] ] + vertices_[ t->v[(i+1)%3] ] + vertices_[ t->v[(i+2)%3] ]
            ++ks[t->v[i]];
        }
    }

    for ( std::vector<Point3D32f>::iterator nv = newverts_.begin (); nv != newverts_.end(); ++nv)
    {
        int num = ks[ nv-newverts_.begin()];
        float beta = 3.0f / (8.0f * (float)num);
        if ( num == 3 )
            beta = 3.0f / 16.0f;

        //float interm = (3.0/8.0 - 0.25 * cos ( 2 * M_PI / num ));
        //float beta = 1.0/num*(5.0/8.0 - interm*interm);

        *nv = (*nv * (0.5f * beta)) + vertices_[nv-newverts_.begin()] * ( 1 - beta*num);
    }
}
```







# Least squares object fitting

Clemson University  
TigerPrints

All Theses

12-2007

Least Squares Fitting of Analytic Primitives on a GPU



$$F(x_0, y_0, z_0, r_0) = \sum \sqrt{(x_i - x_0)^2 + (y_i - y_0)^2 + (z_i - z_0)^2 - r_0^2}^2 \quad (4.34)$$

The partial derivatives of the function  $f$  with respect to each of the parameters

$x_0, y_0, z_0, r_0$  are given by

$$\frac{\partial f}{\partial x_0} = -(x_i - x_0) / (f + r_0) \quad (4.35)$$

$$\frac{\partial f}{\partial y_0} = -(y_i - y_0) / (f + r_0) \quad (4.36)$$





$$\frac{\partial f}{\partial z_0} = -(z_i - z_0) / (f + r_0) \quad (4.37)$$

$$\frac{\partial f}{\partial r_0} = -1 \quad (4.38)$$

# Surface normal from depth images


## Adaptive Neighborhood Selection for Real-Time Surface Normal Estimation from Organized Point Cloud Data Using Integral Images

S. Holzer<sup>1,2</sup> and R. B. Rusu<sup>2</sup> and M. Dixon<sup>2</sup> and S. Gedikli<sup>2</sup> and N. Navab<sup>1</sup>






# As-Rigid-As-Possible Surface Deformations (1-2 students)



[http://www.igl.ethz.ch/projects/ARAP/arap\\_web.pdf](http://www.igl.ethz.ch/projects/ARAP/arap_web.pdf)

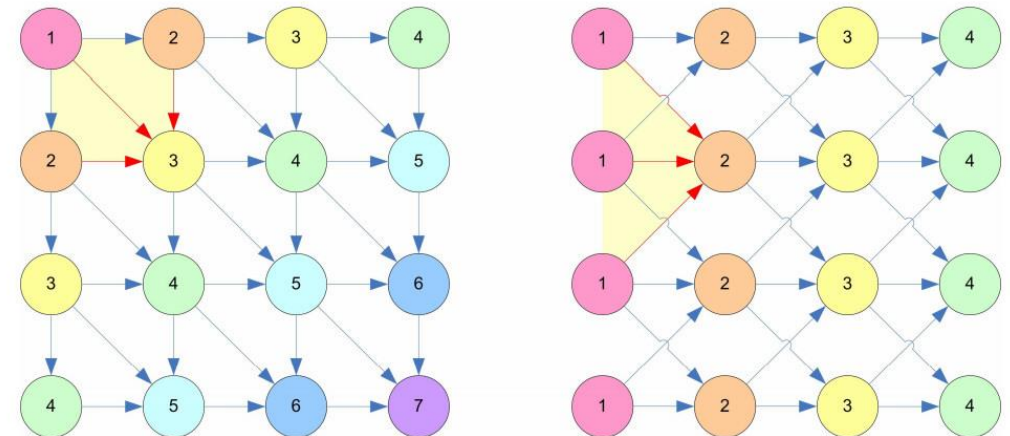



$$E(C_i, C'_i) = \sum_{j \in \mathcal{N}(i)} w_{ij} \|(\mathbf{p}'_i - \mathbf{p}'_j) - \mathbf{R}_i(\mathbf{p}_i - \mathbf{p}_j)\|^2$$

$$\begin{aligned} \sum_j w_{ij} (\mathbf{e}'_{ij} - \mathbf{R}_i \mathbf{e}_{ij})^T (\mathbf{e}'_{ij} - \mathbf{R}_i \mathbf{e}_{ij}) &= \\ = \sum_j w_{ij} (\mathbf{e}_{ij}^T \mathbf{e}'_{ij} - 2 \mathbf{e}_{ij}^T \mathbf{R}_i \mathbf{e}_{ij} + \mathbf{e}_{ij}^T \mathbf{R}_i^T \mathbf{R}_i \mathbf{e}_{ij}) &= \\ = \sum_j w_{ij} (\mathbf{e}_{ij}^T \mathbf{e}'_{ij} - 2 \mathbf{e}_{ij}^T \mathbf{R}_i \mathbf{e}_{ij} + \mathbf{e}_{ij}^T \mathbf{e}_{ij}) &= \end{aligned}$$



C++ source code available  
SVD minimization with CUDA library

# Calculation of Geodesic distances (1-2 students)



# GPU ray tracer (1-2 students)

Eastern Washington University  
**EWU Digital Commons**

---

EWU Masters Thesis Collection

---

2013

## GPU ray tracing with CUDA

Thomas A. Pitkin  
*Eastern Washington University*


<http://www.kevinbeason.com/smallpt/>



This paper uses an iterative ray tracing algorithm outlined by Alejandro Segovia, Xiaoming Li, and Guang Gao (Segovia et al., 2009) to replace the original, recursive ray tracing algorithm. By viewing each recursive ray bounce as an iteration, a “for loop” is implemented into the parallel ray tracer to remove the necessity of recursion without losing functionality.



# CUDA volume raytracer



<http://www.rpenalva.com/blog/?p=229>

CUDA code available, but needs adaptations and profiling

A simple and flexible volume rendering framework for graphics-hardware-based raycasting

4  
Author(s)

S. Stegmaier; Inst. for Visualization & Interactive Syst., Stuttgart Univ., Germany; M. Strengert; T. Klein; T. Ertl

